

Mixture weight and programme strength: construct validation of single-cell topic loadings in diabetic foot skin

Zeyu Fu

Abstract

A higher mean single-cell programme score can reflect either stronger expression within cells or more cells occupying a high state. We examined this distinction before interpreting a fibroblast-associated topic as a diabetic-foot-ulcer (DFU) healing marker. A frozen 15-topic model represented 76,461 cells from 25 selected foot-skin samples. We assessed distributional semantics, technical artifacts, analysis-constant sensitivity, negative controls and external projections. An author-derived mapping collapsed 14 DFU samples to 11 patients; anatomical controls and four representation benchmarks used the corresponding analysis units. Across 25 samples, mean fibroblast loading correlated with the threshold-defined high-state fraction at $\rho = 0.933$ (sample-bootstrap 95% interval 0.789–0.983). The highest fraction, 0.841, occurred in healthy non-diabetic skin. Depth matching preserved 99.5% of state labels, and exclusion of cells carrying immune transcripts preserved sample weights ($\rho = 0.989$). Semantic associations persisted over tested thresholds and topic numbers; the healthy-skin counterexample failed at $K = 10$, and healing-arm significance was threshold-sensitive. The corrected all-cell mean-loading contrast in 11 patients was 0.0280 (permutation $p = 0.2719$; bootstrap interval $[-0.0017, 0.0631]$). The separate 9-versus-5 high-state-fraction design had 24.9% simulated power at its observed effect. Patient-level topic, module, PCA and NMF AUCs were 0.714, 0.786, 0.500 and 0.125; these exploratory estimates do not establish clinical superiority. Paired anatomy remained detectable in 10 patients. External cohorts did not supply independent patient-level healing validation. The mean loading carries information about the prevalence of high-scoring cells, while state-specific means also vary. Its correlation with the threshold fraction does not prove invariant within-state intensity, and the tested fibroblast pattern remains an unvalidated prognostic marker. Patient-level validation requires independently defined outcomes and a clinically justified effect margin.

Keywords: topic model; cell-state abundance; construct validity; DFU; CHI3L1; TIMP1; patient-level inference.

1 Introduction

Diabetic foot ulcers arise in tissue where inflammation, matrix remodeling and vascular impairment can persist together. Even after closure, recurrence remains common [2, 26]. Single-cell studies have described fibroblast-associated programmes that differ between healing groups [19, 24]. Such observations make a programme score a plausible candidate measurement, but they leave a basic question unresolved: what changes when its average is higher in one biopsy than another?

There are at least two contributions. Cells already expressing a programme may express it more strongly, or the biopsy may contain a larger fraction of cells with high scores. Both can occur together. Topic models and gene-expression decompositions provide useful representations of heterogeneous cells [5, 15], and neural latent-variable models can summarize expression for subsequent comparisons [18]. Their coefficients nevertheless acquire biological meaning through the measurement and sampling scheme. A simplex constraint does not make a fitted coefficient an independently validated cell-state probability, and compositional changes can affect a tissue summary without an equivalent change within each cell [1, 6].

This distinction is difficult to resolve from a correlation between a mean and a threshold fraction: both are functions of the same cell scores. Distributional descriptions therefore need to be accompanied by an explicit decomposition of the mean and an examination of the terms treated as approximately constant. Technical explanations require separate attention. Dissociation-induced expression and ambient RNA can alter recovered profiles [25, 27], while normalization and cell selection affect which biological contrasts remain visible [13].

The unit of clinical inference introduces another constraint. Repeated biopsies and thousands of cells from a patient do not supply independent healing outcomes. Sample-aware methods address this dependence when testing expression or abundance changes [9, 17, 22, 28]. Prognostic evaluation additionally requires a defined outcome, a relevant validation population and uncertainty assessed at the patient level [8]. Examining several readouts in a small discovery cohort can inform a later study, but selecting among them on that same evidence risks optimistic performance estimates [7]. Similarly, a nonsignificant difference is not evidence of equivalence without a justified margin [16, 21].


We examine a frozen fibroblast-associated topic in public diabetic-foot skin data, first as a tissue measurement and then as an exploratory healing readout. The analysis distinguishes mean loading within fibroblasts, the fraction above a fitted threshold and the mean over all recovered cells. It combines distributional and technical checks with author-mapped patient comparisons and paired anatomical controls. External wound datasets extend the tissue contexts in which the frozen score can be examined, while their available metadata determine whether they can address clinical validation.

2 Results


2.1 The frozen representation is structured, but its principal loading is a mixture coordinate

The analysis architecture is summarized in Figure 1. The frozen expression map is followed by a cell-state mixture audit, technical and environmental-RNA controls, patient-unit correction, anatomical controls and bounded external construct projections.


A From skin specimens to three distinct readouts



B What changes the average?



C How is the interpretation tested?



Cell-state abundance, technical robustness and healing prediction require separate evidence

Figure 1. Study design and measurement logic. (A) The discovery series supplies 76,461 retained cells from 25 selected foot-skin specimens. A frozen 15-topic expression model produces three distinct summaries: the mean fibroblast loading, the fraction of fibroblasts above a fixed loading cut, and the all-cell mean within each mapped patient. The gene labels describe leading decoder genes, not a newly defined cell type. (B) The cell cartoons are conceptual, not observed counts: grey and teal circles denote low and high states. Changing their relative abundance can increase the sample mean without changing either state's intensity; a mean–fraction correlation alone cannot establish constant within-state expression. (C) Technical perturbations test robustness, the healing contrast collapses 14 DFU specimens into 11 patients, paired anatomy tests tissue sensitivity, and external projections assess cohort-specific construct associations. Each evidence layer addresses a different question; these analyses do not supply independent prognostic validation.

The discovery representation retained 76,461 quality-controlled cells from 25 selected foot-skin samples in the original 54-sample series. Beta perplexity was 243, and the manifold audit passed its prespecified integrity checks. Fibroblast topic was enriched for TIMP1, COL3A1, DCN, FN1, LUM and CHI3L1, with immediate-early and oxidative-stress genes also present among the top-ranked terms. The gene list therefore supports a fibroblast-associated inflammatory/stress description; it does not by itself establish a unique cell state or a healing mechanism (Figure 2).

A pooled two-Gaussian fit supplied a descriptive threshold; this fit did not test a continuous-shift model against a biological mixture model. High-loading samples retained a low-state tail overlapping the low-loading samples. Across the 25 analyzable samples, the sample mean and estimated mixture weight had Spearman $\rho = 0.933$. The bootstrap 95% interval was 0.789–0.983, and leaving out any one sample gave values from 0.924 to 0.957; the permutation null was centred near zero (Table S8). This is the central measurement result: in this dataset, the sample average is strongly associated with the high-state fraction. The two summaries are calculated from the same cells, so their correlation alone does not prove that component intensities are constant. The result does not say that high-state cells have no intrinsic biology. Together with the low-state tail, it shows that the mean contains information about high-score prevalence; it does not identify the relative biological contributions of prevalence and within-state expression. Patient aggregation gave a correlation of 0.937 across 20 patients, with a 10,000-draw patient-bootstrap 95% interval of 0.776–0.990. Omitting one patient at a time gave correlations of 0.927–0.964 (Table S16), extending the specimen-resampling analysis without adding independent observations.

An exact decomposition of all 19,410 discovery fibroblasts clarified the remaining within-state contribution. Empirical low-bin means ranged from 0.0120 to 0.1021, and nonempty high-bin means from 0.2264 to 0.5778. Nine of 25 specimens had an empty high bin. Replacing specimen-specific bin means with their cell-pooled values left a residual RMS of 0.01575 and a residual range of -0.02434 to 0.04377 . The exact mean identity held to below 10^{-15} in every specimen (Table S19). The constant-state approximation is therefore explicit and imperfect; the correlation cannot establish invariant programme intensity.




Figure 2. Association between mean topic loading within fibroblasts and the high-state fraction. (A) Each point is one of 25 GSE165816 specimens: 9 healing (green circles), 5 non-healing (orange triangles) and 11 healthy non-diabetic (purple squares). The high-state fraction uses the pooled cut 0.1842958; Spearman $\rho=0.933$. (B) Sample-shape counts use within-sample loading quantiles: pure low requires the 90th percentile at most 0.15, pure high the 10th percentile at least 0.10, and mixed the 10th percentile at most 0.05 and 90th percentile at least 0.40; other samples form the remaining category. (C) The dot is the observed correlation, the blue segment is the 10,000-draw sample-bootstrap 95% interval (0.789–0.983), and the orange segment is the leave-one-sample-out range (0.924–0.957). Both summaries use the same cells; the association does not by itself prove stable within-state intensity or a healing effect.

2.2 Technical controls support state existence, not prognostic validity

The high state survived the three artifact challenges. The dissociation stress standardized mean difference was +0.246, below the rejection threshold. PTPRC positivity was 0.98% in high-state cells and 0.62% in low-state cells, a rate ratio of 1.60. Exact panel-count matching produced 99.5% cell-label agreement and a sample-weight correlation of 0.984. The state was therefore not explained by a single depth imbalance or a simple doublet excess (Figure 3).




Figure 3. Technical and immune-transcript exclusion controls in the discovery cohort. (A) Dissociation standardized mean difference (SMD, 0.246) and the high-state/low-state PTPRC-positive rate ratio (1.60) are divided by rejection bounds 0.8 and 2.0; the dashed line is one. The rounded state cut is 0.184. (B) Each point is one of 25 samples after restricting to 14,561 fibroblasts eligible for panel-count matching to 986 counts per cell. Reprojection preserves sample fractions ($\rho=0.984$) and 99.5% of cell labels; the dashed line is identity. (C) Bars are BIC for one Gaussian minus BIC for two Gaussians fitted to sample high-state fractions, using 24 matched samples before and after exclusion. Positive values favour two components. Exclusion retains 11,262 of 19,410 fibroblasts overall, and matched sample fractions correlate at $\rho=0.989$. This control measures invariance to the exclusion rule; it does not identify every source of ambient RNA or tissue-composition dependence.

Removing cells carrying any of the prespecified immune transcripts left the sample weights nearly unchanged ($\rho = 0.989$) and increased the separation of the two components. This supports stability under the specified immune-transcript exclusion rule; it does not identify or exclude ambient RNA contamination. The placebo analysis was deliberately retained: unrelated tissue covariates flattened residual bimodality as well as immune covariates. At 24 matched samples, the data cannot determine whether the state is statistically separable from every tissue-composition axis. The matched exclusion comparison contains 24 samples, while each residual regression uses 25 samples. The immune and placebo regressions explained 49.5% and 35.5% of the variance, respectively, and both residual BIC differences favoured one component (Table S11). A clean technical control is not evidence that the state is a healing marker.

An exploratory scan of all 31 available covariates found 14 associations below BH-adjusted $q = 0.05$ across 25 samples. The largest positive association was with the assigned B/plasma-cell fraction ($\rho = 0.828$); CFD/APOD-rich topic mean loading was negatively associated ($\rho = -0.766$). Fibroblast PTPRC positivity and median UMI count were also associated ($\rho = 0.727$ and 0.663). The complete scan, including non-significant covariates, is in Table S10. These marginal, sample-unit correlations do not establish causal drivers or independence from tissue composition; the positive depth association remains relevant despite the successful count-matching control.

2.3 A real state is not the same as a healing marker

The high-state pattern was not restricted to samples labelled as healing. Four of nine healing samples and none of five non-healing samples were classified as high-mode samples, Fisher exact $p = 0.2208$. The most extreme high-state fraction was a healthy non-diabetic foot-skin sample, whose mixture weight was 0.841. That observation is an existence counterexample to the statement that the state is uniquely a marker of a healing DFU. It neither identifies the tissue process responsible nor refutes a possible probabilistic association with healing.

The lineage assignment used markers disjoint from the tested topic: the marker lists shared no genes. Fibroblast composition itself differed between the healing and non-healing sample groups only descriptively (ratio 1.42, $p = 0.438$). Within the fibroblast gate, the discovery contrast was 2.89-fold, with a bootstrap 95% interval of 1.06–13.97 and sample-level $p = 0.04196$. This value is retained as a discovery observation, not as an independent patient validation, because the 14 samples represent 11 patients and include repeated same-wound, same-day specimens. The complete lineage and composition contrasts are in Table S12.

For the separate all-cell mean-loading readout, collapsing repeats to 11 patients changed the healing-minus-non-healing difference from 0.0435 to 0.0280. This is not a patient-level recalculation of the 2.89-fold fibroblast-restricted contrast. The exhaustively enumerated, add-one-corrected permutation p value was 0.2719 and the bootstrap 95% interval was $[-0.0017, 0.0631]$. The earlier 14-sample contrast was 0.0435 with $p = 0.0764$ and sensitivity margin 0.0753; the patient-unit margin was 0.0618. Neither margin is a clinical MCID. The corrected result is therefore a unit-aware discovery result with substantial uncertainty, not a definitive null test. Mapped scores are given in Table S2; equal weighting of repeated specimen means gave a similar patient contrast of 0.0284. The full equivalence sensitivity grid is in Table S15. Patient omission left the all-cell contrast positive, from 0.0139 to 0.0347 across 11 deletions (Table S16); this is an influence range, not a confidence interval or evidence of clinical discrimination.

2.4 The negative healing result has a measured power boundary

The discovery design had only 24.9% estimated power at the observed healing-minus-non-healing effect. Under the stated simulation, the minimum detectable difference at 80% power was 0.519, approximately 1.9 times the observed mean scale. Figure 4C displays this as a power boundary rather than as a null result. This calculation describes the effect size detectable with the specified design; it does not by itself exclude large or moderate healing effects (Figure 4).




Figure 4. State abundance, analysis unit and design power. (A) Symbols show high-state fractions for 9 healing, 5 non-healing and 11 healthy specimens; horizontal segments mark arm means. Green circles, orange triangles and purple squares identify these groups. The largest fraction is 0.841 in healthy non-diabetic foot skin, showing that the state is not exclusive to healing ulcers. (B) Dots and within-arm bootstrap 95% intervals show the healing-minus-non-healing difference in all-cell mean fibroblast-associated topic loading for 14 specimens and 11 mapped patients. The patient analysis resamples 7 healing and 4 non-healing patients in 30,000 draws. Its readout differs from the fibroblast high-state fraction in A; the vertical dashed line is zero. (C) Gaussian location-shift simulation with a two-sided Mann–Whitney test estimates 24.9% power for the separate 9-versus-5 high-state-fraction design, against an 80% target. Its minimum detectable difference is 0.519. These are design sensitivities, not effect-exclusion bounds or clinical minimum important differences.

The paired anatomical control confirmed that the frozen encoder responds to a known tissue difference. Across 10 paired patients, the fibroblast-associated topic increased in foot relative to forearm skin by $+0.0476$ (exhaustive sign-flip $p = 0.0107$, BH $q = 0.0402$); the fibroblast-restricted contrast was $+0.1398$ ($p = 0.0146$). Four of 15 topics remained significant after BH

correction (Table S14). The earlier 3-of-15 sample-level positive-control result is directionally consistent, but the paired analysis is the more appropriate patient-unit sensitivity check. These controls make a complete assay failure less likely, but they do not turn the underpowered DFU comparison into a proof of no association.

We also ran a design calculation for the next cohort, keeping it separate from the biological results. Under the current arm-specific spread, one observation per patient, 80% target power and no missing outcomes, a candidate difference of 0.03 would require 41 total observations, while a candidate equivalence margin of 0.05 would require 17. These values are design inputs rather than findings: the clinical MCID is still open, repeated samples increase the design effect, and the calculation cannot turn the discovery cohort into an adequately powered study. Tables S6 and S7 give the full difference, equivalence and exploratory AUC-gap design calculations.

2.5 Method constants expose both stable and fragile claims

The mean-versus-weight association remained positive across the cut grid, from $\rho = 0.909$ to 0.948 , and across $K=10, 15$ and 20 , from $\rho = 0.924$ to 0.947 . The measurement-semantic claim is therefore not an artefact of the published cut alone (Figure 5).




Figure 5. Threshold, capacity and negative-control analyses. (A,B) Spearman correlation of fibroblast mean loading with the high-state fraction across 25 samples and the two-sided Mann–Whitney healing-arm p value across 9 versus 5 samples, for all 18 cuts from 0.08 to 0.40. Vertical dashed lines mark the fitted cut; the horizontal line in B marks $p=0.05$. (C) Correlations for $K=10, 15$ and 20 after tracking the TIMP1-carrying topic; orange marks $K=10$, where the highest fraction is no longer in healthy skin. (D) The observed high-state-fraction arm difference is compared with the central 95% of 20,000 label permutations ($p=0.1358$). (E) Gene-column shuffling of 6,000 cells raises mean cell perplexity from 3.68 to 14.43 out of 15, violating the concentration criterion. (F) The anatomical sample-level control compares 33 foot and 11 forearm specimens: observed $p=0.0015$, with 5.0% of permuted comparisons below 0.05. This calibration is distinct from the paired 10-patient anatomy analysis and does not establish pipeline-wide error control.

The healthy-skin counterexample was more conditional. It held at all 18 tested cuts, but it failed at $K=10$ because the topic carrying TIMP1 belonged to a different IGF1P-rich programme. It held at $K=15$ and $K=20$ when the CHI3L1/COL3A1/DCN/FN1 state resolved as its own topic. The defensible wording is consequently “the counterexample holds for the tested $K=15$ and $K=20$ fits but not the $K=10$ fit,” not “the counterexample is invariant to model capacity.”

The healing-arm p value moved between 0.042 and 0.062 over the cut grid, with only 4 of 18 cuts below 0.05 (Tables S18 and S17). That sensitivity is a reason to avoid presenting the discovery p value as a stable healing signal. It also independently supports the decision to keep the healing claim unvalidated.

2.6 Negative controls and representation comparisons constrain interpretation

The label-permutation negative control produced an observed difference of +0.2380 inside the 95% null interval $[-0.3183, +0.2874]$, with $p = 0.1358$. The observed contrast was not exceptional under this label-exchangeability null. This single cali-

bration exercise does not prove that the complete pipeline is incapable of false-positive associations. In the unstructured-input control, real data had perplexity 3.68/15, whereas gene-column shuffling produced 14.43/15 and a nearly uniform configuration that failed the integrity criterion; the destroyed co-expression structure was therefore rejected. In the permutation positive control, the observed p value was 0.0015 and the proportion of significant permutations was 5.0%, consistent with the nominal type-I error rate (Figure 5D–F).

The representation benchmark gave point AUCs of 0.822 for fibroblast-associated topic, 0.844 for the module score, 0.533 for PCA and 0.878 for NMF. No pairwise bootstrap interval for an AUC difference excluded zero. The module score alone separated its permutation null at $p = 0.0424$; fibroblast-associated topic and NMF were just above the threshold at $p = 0.0604$ and $p = 0.0597$, while PCA gave $p = 0.4527$. The null AUC means were not all 0.5 (the fibroblast-associated topic, 0.561; NMF, 0.553), illustrating why choosing orientation inside each fold can inflate a naive 0.5 comparison. These results quantify differences without supporting a ranking (Figure 6; Tables S3, S5 and S13).




Figure 6. Sample-unit representation comparisons in 14 DFU specimens (9 healing, 5 non-healing). (A) Points are leave-one-sample-out areas under the ROC curve (AUCs); horizontal bars are percentile 95% intervals from 4,000 arm-stratified bootstrap draws of fixed out-of-fold scores, with common resampling indices across methods. Orange ticks mark the permutation-null means after rerunning outcome-dependent fitting and orientation. The dashed line at AUC=0.5 is a visual reference, not the fitted permutation null. Fibroblast topic denotes the frozen all-cell mean, Module the specified fibro-inflammatory score, PCA principal component analysis, and NMF nonnegative matrix factorization. No paired AUC-difference interval excludes zero. (B) Cells contain Spearman correlations of the 14 sample-level readouts; the colour scale and printed values encode the same statistic. These exploratory comparisons do not establish a clinical method ranking.

2.7 Patient-level readouts weaken outcome discrimination, while anatomy remains detectable

The patient-unit benchmark lowered the apparent discrimination of the frozen topic readout relative to a sample-level analysis and made the uncertainty explicit. Across 11 patients, fibroblast-associated topic AUC was 0.714 (bootstrap 95% interval 0.357–1.000; permutation $p = 0.302$), while the module score was 0.786 (0.500–1.000; $p = 0.166$). PCA was 0.500 and NMF was 0.125, with broad intervals and p values of 0.503 and 0.975. The companion sample-level score matrix shows correlations from 0.165 to 0.723 across topic, module, PCA and NMF readouts, so relatedness does not imply a stable ordering. Topic versus module and topic versus PCA differences were not separated from zero; only the topic-versus-NMF and module-versus-NMF point contrasts had bootstrap intervals excluding zero. Because this is a single, small discovery cohort, the result is a unit-sensitivity audit rather than a claim that one representation is clinically superior (Tables S4 and S13; Figure 7).

The paired anatomical control provides the complementary positive result. The same fibroblast-associated topic readout changed between foot and forearm skin in 10 paired patients, with exact $p = 0.0107$ and BH $q = 0.0402$. The combination of weak, uncertain outcome discrimination and a detectable anatomical contrast supports a measurement-sensitivity interpretation, but it does not establish healing specificity.

2.8 External projections remain construct audits

The external immune-axis analysis used mutually exclusive lineage compartments. GSE223964 and GSE245703 gave immune-axis correlations of -0.024 and $+0.846$, respectively. Within the immune compartment, B/plasma-cell share was weakly related to the topic signal (-0.071 and $+0.077$), while correlations with median detected genes per cell were negative in both datasets (-0.429 and -0.427). The signs are inconsistent for a generalized immune-axis effect, and neither dataset has an authoritative patient-level healing endpoint (Table S9; Figure 8).

GSE231643 therefore provides a weaker, label-only replication of direction: the within-fibroblast ratio falls from 2.89 in the discovery analysis to 1.66 (sample-bootstrap 95% interval 1.01–6.10; two-sided Mann–Whitney $p = 0.1429$; Table S12). The separate GSE241132 acute-wound cohort gave a broad immune-axis rho of -0.070 and an author-annotated B/plasma-cell rho of $+0.210$ across three donors; both values are descriptive and were not combined with the DFU outcome analyses.




Figure 7. Patient-unit discrimination and paired anatomical sensitivity. (A) Points are leave-one-patient-out AUCs for 7 healing and 4 non-healing patients after cell-count-weighted collapse of repeated specimens. Bars are 95% intervals from 4,000 arm-stratified patient bootstrap draws of fixed out-of-fold scores; the dashed line marks AUC=0.5. Fibroblast topic, Module, PCA and NMF denote the frozen topic mean, fibro-inflammatory module, principal component and nonnegative matrix factorization readouts. (B) The point is the mean all-cell fibroblast-associated topic loading difference, foot minus forearm, across 10 paired patients; the bar is a 4,000-draw paired-patient bootstrap 95% interval and the horizontal line is zero. The exhaustive sign-flip test with add-one correction gives $p=0.0107$ and Benjamini-Hochberg $q=0.0402$ across 15 topics. Both panels use the discovery source and are not independent clinical validation.




Figure 8. Frozen external projections under mutually exclusive lineage gates. (A) Each point is a cohort-level Spearman correlation between the sample fibroblast fibroblast-associated topic mean and either immune fraction among all QC cells (teal circles) or B/plasma share within immune cells (purple squares). (B) Orange points correlate the same topic mean with median detected genes per cell. GSE223964, GSE245703 and GSE268834 contribute 8, 12 and 8 specimens; vertical dashed lines mark zero. (C) Points show diabetic-minus-non-diabetic differences in the sample fibroblast mean: +0.0464 in GSE268834 (descriptive permutation $p=0.2812$) and -0.0204 in GSE248247 ($p=1.0$). GSE248247 contains one specimen per group and 358 QC cells, so its correlations are not estimated. No panel displays confidence intervals. These cohorts lack qualifying patient-level healing endpoints and are not pooled for prognosis inference.

GSE268834 supplied another descriptive projection: the topic-associated fibroblast signal correlated with exclusive immune fraction at +0.714, and the diabetic-minus-non-diabetic fibroblast-topic contrast was +0.0464 with $p = 0.2812$. This is compatible with an immune/tissue-context hypothesis but cannot distinguish it from other between-sample differences.

GSE248247 was a deliberately low-yield stress test of the frozen projection. It contained 561 raw cells and 358 quality-controlled cells, with 6,994 of 7,002 panel genes covered in each sample. The fibroblast topic contrast between one diabetic and one non-diabetic sample was -0.0204 with descriptive permutation $p = 1.0$. All correlations were non-estimable because $n=2$. The exclusive compartment sum was exactly one in both samples, demonstrating that the lineage gate was non-overlapping; the dataset cannot validate healing or support pooled patient inference.

3 Discussion

Mean fibroblast loading and the fraction of high-scoring fibroblasts were closely associated across the discovery specimens, and the association persisted after patient aggregation. This establishes that the mean conveys information about the prevalence of high scores in the recovered cells. It does not make the two measurements interchangeable. The low- and high-bin means varied across

specimens, and nine specimens contained no cells in the high bin. A constant-state approximation retained a nonzero specimen-level residual. The appropriate interpretation is therefore a mixture of prevalence and within-bin score variation, conditional on the chosen representation and threshold.

This reading is consistent with the distinction between cell identity, cellular activity and tissue composition [6, 15]. A reproducible expression loading need not identify a discrete biological state, particularly when a continuous bounded score is summarized by a two-Gaussian fit. The threshold provides an operational measurement, not evidence that fibroblasts occupy two stable attractors. Spatial measurements of cell abundance, independent protein or transcript assays and repeated tissue sampling would help determine how much of the observed variation reflects tissue organization, cell activity or recovery during processing.

The technical controls narrow some explanations without resolving all of them. Most state labels survived count matching, and specimen summaries were stable after excluding cells carrying specified immune transcripts. Those observations support robustness to the tested perturbations. They cannot establish the absence of ambient RNA or doublets, whose effects depend on the source of contamination and the expression profiles of affected cells [25, 27]. The similar attenuation of residual bimodality by immune and unrelated tissue covariates further limits an interpretation specific to one immune mechanism.

The clinical evidence is more limited than the measurement association. The highest threshold fraction occurred in healthy skin, challenging disease exclusivity without excluding every possible prognostic relationship. The patient-level healing contrast was imprecise, and the modest difference between exhaustive and add-one permutation probabilities did not change its interpretation. Paired anatomical differences were detectable, but tissue sensitivity alone does not establish healing specificity. These distinctions matter when cell-level observations are translated into claims about patients [8, 22, 28].

The representation comparison is an internal audit of this cohort. Although outcome-dependent orientation and component selection were performed within training folds, the frozen topic encoder had already seen the discovery expression data. Bootstrap intervals on fixed out-of-fold scores omit uncertainty from repeating the full learning procedure, and the folds share training patients [3]. Differences in observed AUC therefore do not establish a stable ranking of methods. Searching across representations, thresholds and topic capacities also requires independent evaluation before a preferred readout can be treated as a validated predictor [7].

External projections demonstrate that the frozen readout can be computed in additional wound contexts, but the cohorts do not share an independently adjudicated patient-level healing endpoint. Variation in sampling, cell yield and compartment definitions can contribute to inconsistent immune associations. These datasets cannot be pooled as interchangeable replications of a prognostic claim. The smallest projection, with two specimens, is informative about computational feasibility rather than association strength.

A future validation study should specify the patient outcome and clinically meaningful difference before examining scores, preserve the patient as the independent unit and freeze the tissue-processing and analysis rules. The present power calculations depend on estimated discovery variances and approximate score distributions; they do not determine a guaranteed recruitment target. Likewise, the smallest margin compatible with an equivalence test is a description of these data, not a clinical threshold [16]. Independent patient-level evaluation is needed to establish whether an interpretable tissue readout improves prognosis beyond the information already available at clinical assessment.

4 Methods

4.1 Cohort selection and independent units

GSE165816 comprises 54 original samples and 27 author-mapped participants [19, 24]. The frozen model used 25 selected foot-skin samples in the non-diabetic, DFU-healing and DFU-non-healing groups: 81,602 cells before cell quality control and 76,461 afterwards. Cells required at least 200 detected genes and a mitochondrial count fraction at most 0.20; genes required detection in at least 10 cells before panel selection. The final panel contained 7,002 genes, obtained from a nominal 7,000-gene dispersion selection with retained anchor genes. The discovery outcome subset contained 9 healing and 5 non-healing samples.

The clinical supplement maps those 14 outcome samples to 11 patients: 7 healing and 4 non-healing. G4/G23, G7/G8 and G33/G34 are repeated samples from the same wound on the same day. We combine each patient’s retained cells before computing the primary patient score. The outcome is the study’s healing category within a 12-week follow-up window, not a measured individual healing time. Series metadata describe a shared Emory 10x Chromium 3’ V2/V3 workflow and NovaSeq instrumentation; the absence of per-sample chemistry and flow-cell identifiers prevents run-specific batch adjustment.

4.2 What the frozen neural topic model computes

For nonnegative panel counts x_{ig} with positive total, let $u_{ig} = x_{ig} / \sum_{h=1}^G x_{ih}$. The implemented model uses a diagonal-Gaussian encoder in unconstrained coordinates and a softmax decoder:

$$q_{\phi}(z | u_i) = \mathcal{N}(\mu_{\phi}(u_i), \text{diag } \sigma_{\phi}^2(u_i)), \quad (1)$$

$$\theta_{ik} = \frac{e^{z_{ik}}}{\sum_{\ell=0}^{K-1} e^{z_{i\ell}}}, \quad \beta_{kg} = \frac{e^{b_{kg}}}{\sum_{h=1}^G e^{b_{kh}}}, \quad \hat{u}_{ig} = \sum_{k=0}^{K-1} \theta_{ik} \beta_{kg}. \quad (2)$$

Equations (1)–(2) define the encoder, simplex coordinates and normalized topic-gene decoder used below. Thus $\theta_i \in \Delta^{K-1} = \{a \in \mathbb{R}^K : a_k \geq 0, \sum_{k=0}^{K-1} a_k = 1\}$ and every topic-gene row sums to one. The topic indices start at zero. A Gaussian mapped through softmax gives a logistic-normal representation [1]. This differs from a Dirichlet prior, while retaining the mixture representation that motivates topic models [5, 23]. Training minimizes

$$\mathcal{L}(\phi, b) = \frac{1}{N} \sum_i \mathbb{E}_{q_\phi} \left[- \sum_g u_{ig} \log(\hat{u}_{ig} + 10^{-10}) \right] + 0.01 \frac{1}{N} \sum_i \text{KL}(q_\phi(z | u_i) \| \mathcal{N}(0, I)). \quad (3)$$

Equation (3) is the fitted objective: each cell contributes cross-entropy on its panel proportions, followed by a Gaussian regularizer. It is neither a depth-weighted count likelihood nor an unmodified evidence lower bound for such a likelihood. With encoder log-variance $\ell_{ik} = \log \sigma_{ik}^2$, the reparameterization and regularizer are

$$z_{ik} = \mu_{ik} + e^{\ell_{ik}/2} \epsilon_{ik}, \quad \epsilon_i \sim \mathcal{N}(0, I), \quad \text{KL}_i = \frac{1}{2} \sum_k (\mu_{ik}^2 + e^{\ell_{ik}} - 1 - \ell_{ik}). \quad (4)$$

Equation (4) is the diagonal-Gaussian calculation used by the implementation [14]; one noise draw per cell estimates the reconstruction expectation in each training batch. The encoder has widths 256 and 128, batch normalization, ReLU activations and dropout 0.1. The recorded fit used $K = 15$, 40,000 Adam steps, learning rate 10^{-3} and batches of 512. Deterministic projection disables encoder dropout and uses $z = \mu_\phi(u)$ and $\theta = \text{softmax}(\mu_\phi(u))$, which is not generally $\mathbb{E}_q[\text{softmax}(z)]$. Missing panel genes in external projections are set to zero. Normalization and quality control can alter biological interpretations, so this definition is kept distinct from other variance-stabilizing methods [11, 13].

Topic-gene perplexity is summarized as $\exp\{-K^{-1} \sum_{k,g} \beta_{kg} \log \beta_{kg}\}$; it measures concentration of decoder weights, not superiority over other representations. The latent coordinate $\mu \in \mathbb{R}^K$ is not the simplex coordinate θ . Softmax is unchanged by adding a common scalar to all coordinates of a cell; reconstruction therefore does not identify that offset. The Gaussian penalty constrains the fitted coordinates but does not give them a unique biological interpretation. Topic weights are decoder coefficients, not independently calibrated probabilities of membership in a biological cell type.

4.3 Threshold fractions and the exact mean decomposition

Let $Y_{si} = \theta_{si,*}$ for fibroblast-gated cell i in sample s . Lineage markers are disjoint from the tested topic-gene list by an explicit overlap check; the gate is nevertheless an expression-based operational definition, not an independently measured cell identity. A pooled two-Gaussian fit provides a descriptive cut $c = 0.1842958$, the midpoint of the fitted component means. Gaussian components approximate a bounded score distribution and are not assumed to be literal biological populations.

$$\bar{Y}_s = \frac{1}{n_s} \sum_i Y_{si}, \quad \hat{w}_s(c) = \frac{1}{n_s} \sum_i \mathbf{1}\{Y_{si} > c\}. \quad (5)$$

Equation (5) defines both the sample mean and the reported hard-threshold fraction; the latter is not an average posterior responsibility of the Gaussian mixture. Define $\bar{Y}_{s,L}$ and $\bar{Y}_{s,H}$ as means below/at and above the cut. Then, with empty-component contributions interpreted as zero,

$$\bar{Y}_s = (1 - \hat{w}_s) \bar{Y}_{s,L} + \hat{w}_s \bar{Y}_{s,H}. \quad (6)$$

Equation (6) is exact. Only under an additional approximate component-stability assumption, $\bar{Y}_{s,L} \approx m_L$ and $\bar{Y}_{s,H} \approx m_H$, does it reduce to $\bar{Y}_s \approx m_L + (m_H - m_L) \hat{w}_s$. Consequently, a large correlation between the two summaries cannot independently prove stable intensity: both summaries use the same cells. Persistence of overlapping low-state tails supplies complementary distributional evidence. The 10th percentile of a high-loading sample (0.029) overlaps the scale of the 90th percentile in a low-loading sample (0.119); this observation alone is not a formal test of multimodality.

To quantify the term omitted by a constant-state approximation, let m_L, m_H be the cell-pooled means in the two threshold bins across the 19,410 discovery fibroblasts. They are empirical bin means, not the fitted Gaussian means. For each specimen,

$$A_s = (1 - \hat{w}_s) m_L + \hat{w}_s m_H, \\ R_s = (1 - \hat{w}_s) (\bar{Y}_{s,L} - m_L) + \hat{w}_s (\bar{Y}_{s,H} - m_H), \quad \bar{Y}_s = A_s + R_s. \quad (7)$$

Equation (7) separates a specified constant-state approximation from its exact residual. An empty bin contributes zero and its undefined mean is omitted from the corresponding mean range. We report specimen-weighted residual RMS, bin-mean ranges and all 25 specimen summaries. These are finite-data descriptions without a test of state invariance or a causal allocation of variance (Table S19).

The semantic statistic is $\rho_{\text{sem}} = \text{cor}_{\text{Spearman}}(\bar{Y}_s, \hat{w}_s)$ across 25 samples. A 10,000-draw sample-row percentile bootstrap, 10,000 row permutations and leave-one-sample-out recalculation characterize its stability. These do not create independent patients. The patient-collapsed correlation of 0.937 across 20 mapped patients is an additional descriptive unit check.

A separate patient-resampling analysis collapses the 25 specimen means and high-state fractions to 20 patients using fibroblast counts as within-patient weights. It draws 10,000 patient samples with replacement (seed 217) and reports the percentile interval for their Spearman correlation. Leave-one-patient-out analyses recompute this correlation across the remaining 19 patients and separately recompute the all-cell healing contrast across the remaining 10 outcome patients. The reported omission range is a sensitivity diagnostic, not an interval estimate (Table S16).

4.4 Technical controls and constant sensitivity

Dissociation-induced programmes and ambient transcripts motivate the technical controls [25, 27]. The stress comparison uses a standardized mean difference (SMD) with rejection bound 0.8; the doublet-oriented control compares PTPRC-positive rates with ratio bound 2.0. Exact panel-count matching reprojects downsampled counts, comparing cell labels and sample fractions. Removing any cell carrying a specified immune transcript tests invariance to that exclusion rule. This cannot by itself identify ambient contamination or exclude every doublet. Residual mixture checks with immune covariates are compared with four unrelated placebo covariates; similarly flattened residuals limit claims of statistical separability.

Sensitivity fits use 18 cuts between 0.08 and 0.40 (including the rounded fitted cut) and $K \in \{10, 15, 20\}$. The TIMP1-carrying topic is tracked descriptively across capacities; topic labels are not assumed exchangeable or biologically identical across K . Label permutation uses 20,000 shuffles. Gene-column shuffling preserves each gene's marginal counts while destroying its joint cellwise association with other genes. The resulting topic concentration and integrity statistics are compared with real input. The anatomical label-permutation calibration checks a finite empirical false-positive fraction; it is not a proof of pipeline-wide error control. These analyses are retrospective computational controls, not a registered clinical trial.

The exploratory covariate scan retains the original rounded rule $Y_{si} \geq 0.184$, used by the technical-control analyses, rather than substituting the fitted cut 0.1842958. It correlates the sample high-state fraction with 14 other fibroblast topic means, 11 whole-sample lineage fractions, total and fibroblast cell counts, fibroblast dissociation score, median UMI count, PTPRC-positive fraction and the numeric G specimen label. Two-sided Spearman probabilities are adjusted across all 31 comparisons by Benjamini–Hochberg. The 25 rows are specimens, include repeated patients, and receive no clinical-outcome adjustment. The G label is an identifier covariate, not a verified sequencing batch or collection time. Topic coordinates and lineage fractions have compositional dependence.

4.5 Patient contrast, permutation and equivalence sensitivity

Let $\mathcal{C}_p$ contain all retained cells in the mapped DFU foot specimens of patient p . The primary corrected score is $S_p = |\mathcal{C}_p|^{-1} \sum_{i \in \mathcal{C}_p} \theta_{i*}$, so repeated specimens contribute in proportion to retained cell counts. This all-cell readout differs from $\bar{Y}_s$ restricted to fibroblasts and from $\hat{w}_s$. For healing set H and non-healing set N ,

$$\hat{\Delta} = \frac{1}{|H|} \sum_{p \in H} S_p - \frac{1}{|N|} \sum_{p \in N} S_p. \quad (8)$$

Equation (8) is the primary patient-unit contrast; it is separate from the sample-level fibroblast ratio and from the threshold fraction. All $\binom{11}{7} = 330$ patient-label allocations are enumerated. To reproduce the recorded result, the implementation retains an add-one correction even for exhaustive enumeration:

$$p_+ = \frac{1 + \sum_{\pi \in \Pi} \mathbf{1}\{|\hat{\Delta}_\pi| \geq |\hat{\Delta}|\}}{1 + |\Pi|}. \quad (9)$$

Equation (9) is the recorded add-one-corrected two-sided permutation probability. It is slightly conservative relative to the uncorrected exhaustive fraction; the add-one convention is usually motivated for randomly drawn permutations [20], whereas all allocations are enumerated here. Exchangeability is an assumption for an observational label comparison, not a randomization supplied by study design. A 30,000-draw bootstrap resamples patients within outcome arms [10]. Small group sizes and conditional resampling explain why bootstrap intervals and permutation tests need not be dual.

An independent enumeration also reports the exhaustive probability without the extra correction:

$$p_{\text{exact}} = \frac{\#\{\pi \in \Pi : |\hat{\Delta}_\pi| \geq |\hat{\Delta}|\}}{|\Pi|} = \frac{89}{330} = 0.2697. \quad (10)$$

Equation (10) includes the observed allocation and numerical ties. It distinguishes exhaustive enumeration from randomly sampled permutations [12]; the recorded corrected value $90/331 = 0.2719$ is retained alongside it. Neither convention makes observational healing labels randomized or tests a causal effect. Table S20 reports the independent arithmetic check, including the Welch quantities below.

For a symmetric equivalence margin M , the Welch two-one-sided-test procedure [16, 21] uses $SE = (s_H^2/n_H + s_N^2/n_N)^{1/2}$ and Welch degrees of freedom ν :

$$p_L = 1 - F_{t_\nu} \left(\frac{\hat{\Delta} + M}{SE} \right), \quad p_U = F_{t_\nu} \left(\frac{\hat{\Delta} - M}{SE} \right). \quad (11)$$

Equation (11) gives the two one-sided tests used for a prespecified margin. Both must be below 0.05 to reject non-equivalence. Writing $\nu_H = s_H^2/n_H$ and $\nu_N = s_N^2/n_N$ makes the degrees of freedom and margin boundary explicit:

$$\nu = \frac{(\nu_H + \nu_N)^2}{\nu_H^2/(n_H - 1) + \nu_N^2/(n_N - 1)}, \quad M_* = |\hat{\Delta}| + t_{\nu, 0.95} SE. \quad (12)$$

Equation (12) is the infimum margin containing the two-sided 90% Welch interval; strict rejection requires $M > M_*$. Independent calculation gives $SE = 0.01823$, $\nu = 8.099$ and $M_* = 0.06181$ for the 11-patient comparison. The original sample-unit boundary is 0.0753. A margin selected after seeing these data is a sensitivity summary, not a clinically justified equivalence claim. No clinical minimum important difference has been supplied.

4.6 Discrimination, anatomy and design simulations

For held-out prediction scores r , AUC is the empirical probability that a healing score exceeds a non-healing score, with half credit for ties:

$$\widehat{\text{AUC}} = \frac{1}{n_H n_N} \sum_{h \in H, n \in N} \left[\mathbf{1}\{r_h > r_n\} + \frac{1}{2} \mathbf{1}\{r_h = r_n\} \right]. \quad (13)$$

Equation (13) is the pairwise empirical AUC used for the held-out scores. Topic orientation, module preprocessing and the choice of PCA/NMF components are determined inside each training fold. The topic encoder itself was trained without outcomes on the discovery cohort, including held-out patients' expression: this is an internal discrimination audit of a frozen discovery representation, not fully external representation validation. Bootstrap intervals resample fixed out-of-fold scores with common arm-stratified indices across methods; the entire fitting/orientation process is rerun for label permutations. All allocations are used for topic and module readouts; PCA and NMF use 200 sampled relabellings. Sample-level comparisons use 14 specimens and patient-level comparisons use 11 mapped patients; they are not interchangeable.

PCA scaling and basis fitting, and NMF fitting, use training rows only. The pooled out-of-fold AUC compares scores from different fitted folds; those scores are not calibrated clinical risks. Reusing the same cohort to inspect readouts and analysis choices limits confirmatory interpretation [7]. Overlapping training sets and conditioning bootstrap draws on fixed scores also leave learning-algorithm uncertainty outside the reported intervals [3].

The paired anatomical control tests foot-minus-forearm differences in 10 mapped patients using all 2^{10} sign allocations and the same add-one convention. Across 15 topics, Benjamini–Hochberg adjustment is $q_{(i)} = \min\{1, \min_{j \geq i} 15p_{(j)}/j\}$, with the usual assumptions required for false-discovery-rate interpretation [4].

The historical high-state-fraction power calculation uses 9 versus 5 specimens, Gaussian location shifts with pooled SD 0.283 and a two-sided Mann–Whitney test. Its 24.9% power and $\text{MDE}_{0.8} = 0.519$ apply to this approximate design, not to the 11-patient all-cell contrast. Gaussian simulation can leave the unit interval, so it is a design-scale sensitivity approximation. A separate next-cohort simulation uses arm-specific SDs of all-cell mean scores, 0.04845 and 0.01584, Welch testing or Welch equivalence testing, and 4,000 draws per candidate size. The candidate totals are approximate conditional calculations, not validated recruitment targets (Table S6).

4.7 External projections and limits of identification

Every external cohort is projected into the frozen panel/model without outcome-guided retuning. The immune and B/plasma gates exclude promiscuous markers such as CD74, CD37, CD38 and SDC1 and assign non-overlapping compartments. Immune fraction uses all retained cells as its denominator; B/plasma share within immune cells uses only that compartment. Their correlations estimate different quantities. Negative depth correlations cannot by themselves exclude depth confounding, because a marginal sign is not a conditional adjustment.

GSE231643 supplies eight title-labelled specimens but lacks an authoritative patient mapping. GSE241132 contributes an acute-wound, three-donor construct summary. GSE223964, GSE245703 and GSE268834 lack qualifying patient-level healing endpoints, and GSE248247 has only two samples. We do not pool these correlations or describe them as independent prognostic validation. No analysis identifies a treatment effect, a time-to-healing model or a causal effect of immune infiltration.

Data and reproducibility

The discovery series is GSE165816. External projections use GSE231643, GSE241132, GSE223964, GSE245703, GSE268834 and GSE248247; each accession is linked in Table S1. This is a secondary analysis of publicly released data, with no new recruitment or intervention. The original studies govern participant consent and source-study ethical approvals. Shared series-level laboratory metadata do not identify sample-specific sequencing batches. The mathematical definitions, analysis units, numerical summaries and figure source data specify the present analysis; no unavailable patient identifiers, healing times or independent clinical validation are imputed. Analysis code is released under the MIT License. Manuscripts, figures and source-study data remain under their respective rights and repository terms. The accompanying software archive, authored by Zeyu Fu, is available at <https://github.com/PeterPonyu/wound-measurement-construct>. An independently archived software release is available at <https://doi.org/10.5281/zenodo.22875656>.

References

- [1] J. Aitchison and S. M. Shen. Logistic-normal distributions: Some properties and uses. *Biometrika*, 67(2):261–272, 1980. doi: 10.1093/biomet/67.2.261.
- [2] D. G. Armstrong, A. J. Boulton, and S. A. Bus. Diabetic Foot Ulcers and Their Recurrence. *New England Journal of Medicine*, 376(24):2367–2375, 2017. doi: 10.1056/nejmra1615439.
- [3] Y. Bengio and Y. Grandvalet. No Unbiased Estimator of the Variance of K-Fold Cross-Validation. *Journal of Machine Learning Research*, 5:1089–1105, 2004. URL <https://www.jmlr.org/papers/v5/grandvalet04a.html>.
- [4] Y. Benjamini and Y. Hochberg. Controlling the False Discovery Rate: A Practical and Powerful Approach to Multiple Testing. *Journal of the Royal Statistical Society: Series B*, 57(1):289–300, 1995. doi: 10.1111/j.2517-6161.1995.tb02031.x.

- [5] D. M. Blei, A. Y. Ng, and M. I. Jordan. Latent Dirichlet Allocation. *Journal of Machine Learning Research*, 3:993–1022, 2003. URL <https://www.jmlr.org/papers/v3/blei03a.html>.
- [6] M. Büttner, J. Ostner, C. L. Müller, F. J. Theis, and B. Schubert. scCODA is a Bayesian model for compositional single-cell data analysis. *Nature Communications*, 12(1):6876, 2021. doi: 10.1038/s41467-021-27150-6.
- [7] G. C. Cawley and N. L. C. Talbot. On Over-fitting in Model Selection and Subsequent Selection Bias in Performance Evaluation. *Journal of Machine Learning Research*, 11:2079–2107, 2010. URL <https://www.jmlr.org/papers/v11/cawley10a.html>.
- [8] G. S. Collins, J. B. Reitsma, D. G. Altman, and K. G. Moons. Transparent Reporting of a multivariable prediction model for Individual Prognosis Or Diagnosis (TRIPOD): The TRIPOD Statement. *Annals of Internal Medicine*, 162(1):55–63, 2015. doi: 10.7326/m14-0697.
- [9] H. L. Crowell, C. Soneson, P.-L. Germain, D. Calini, L. Collin, C. Raposo, D. Malhotra, and M. D. Robinson. muscat detects subpopulation-specific state transitions from multi-sample multi-condition single-cell transcriptomics data. *Nature Communications*, 11:6077, 2020. doi: 10.1038/s41467-020-19894-4.
- [10] B. Efron and R. J. Tibshirani. *An Introduction to the Bootstrap*. Chapman and Hall, New York, 1993. ISBN 9780412042317.
- [11] C. Hafemeister and R. Satija. Normalization and variance stabilization of single-cell RNA-seq data using regularized negative binomial regression. *Genome Biology*, 20:296, 2019. doi: 10.1186/s13059-019-1874-1.
- [12] J. Hemerik and J. Goeman. Exact testing with random permutations. *TEST*, 27(4):811–825, 2018. doi: 10.1007/s11749-017-0571-1.
- [13] L. Heumos, A. C. Schaar, C. Lance, A. Litinetskaya, F. Drost, L. Zappia, M. D. Lücken, D. C. Strobl, J. Henao, F. Curion, Single-cell Best Practices Consortium, H. Aliee, M. Ansari, P. Badia-i Mompel, M. Büttner, E. Dann, D. Dimitrov, L. Dony, A. Frishberg, D. He, S. Hediye-zadeh, L. Hetzel, I. L. Ibarra, M. G. Jones, M. Lotfollahi, L. D. Martens, C. L. Müller, M. Nitzan, J. Ostner, G. Palla, R. Patro, Z. Piran, C. Ramírez-Suástegui, J. Saez-Rodriguez, H. Sarkar, B. Schubert, L. Sikkema, A. Srivastava, J. Tanevski, I. Virshup, P. Weiler, H. B. Schiller, and F. J. Theis. Best practices for single-cell analysis across modalities. *Nature Reviews Genetics*, 24(8):550–572, 2023. doi: 10.1038/s41576-023-00586-w.
- [14] D. P. Kingma and M. Welling. Auto-Encoding Variational Bayes. In *International Conference on Learning Representations*, 2014. URL <https://arxiv.org/abs/1312.6114>.
- [15] D. Kotliar, A. Veres, M. A. Nagy, S. Tabrizi, E. Hodis, D. A. Melton, and P. C. Sabeti. Identifying gene expression programs of cell-type identity and cellular activity with single-cell RNA-Seq. *eLife*, 8:e43803, 2019. doi: 10.7554/eLife.43803.
- [16] D. Lakens. Equivalence Tests: A Practical Primer for *t* Tests, Correlations, and Meta-Analyses. *Social Psychological and Personality Science*, 8(4):355–362, 2017. doi: 10.1177/1948550617697177.
- [17] S. E. Lazic. The problem of pseudoreplication in neuroscientific studies: is it affecting your analysis? *BMC Neuroscience*, 11:5, 2010. doi: 10.1186/1471-2202-11-5.
- [18] R. Lopez, J. Regier, M. B. Cole, M. I. Jordan, and N. Yosef. Deep generative modeling for single-cell transcriptomics. *Nature Methods*, 15(12):1053–1058, 2018. doi: 10.1038/s41592-018-0229-2.
- [19] NCBI Gene Expression Omnibus. GSE165816: Single-cell transcriptomic landscape of diabetic foot ulcers. Gene Expression Omnibus accession, 2022. URL <https://www.ncbi.nlm.nih.gov/geo/query/acc.cgi?acc=GSE165816>.
- [20] B. Phipson and G. K. Smyth. Permutation P-values Should Never Be Zero: Calculating Exact P-values When Permutations Are Randomly Drawn. *Statistical Applications in Genetics and Molecular Biology*, 9(1), 2010. doi: 10.2202/1544-6115.1585.
- [21] D. J. Schuirmann. A comparison of the two one-sided tests procedure and the power approach for assessing the equivalence of average bioavailability. *Journal of Pharmacokinetics and Biopharmaceutics*, 15(6):657–680, 1987. doi: 10.1007/BF01068419.
- [22] J. W. Squair, M. Gautier, C. Kathe, M. A. Anderson, N. D. James, T. H. Hutson, R. Hudelle, T. Qaiser, K. J. E. Matson, Q. Barraud, A. J. Levine, G. La Manno, M. A. Skinnider, and G. Courtine. Confronting false discoveries in single-cell differential expression. *Nature Communications*, 12(1):5692, 2021. doi: 10.1038/s41467-021-25960-2.
- [23] A. Srivastava and C. Sutton. Autoencoding Variational Inference For Topic Models. In *International Conference on Learning Representations*, 2017. URL <https://arxiv.org/abs/1703.01488>.
- [24] G. Theodoridis, B. E. Thomas, D. Sarkar, H. L. Mumme, W. J. R. Pilcher, B. Dwivedi, T. Sandoval-Schaefer, R. F. Sirbulescu, A. Kafanas, I. Mezghani, P. Wang, A. Lobao, I. S. Vlachos, B. Dash, H. C. Hsia, V. Horsley, S. S. Bhasin, A. Veves, and M. Bhasin. Single cell transcriptomic landscape of diabetic foot ulcers. *Nature Communications*, 13:181, 2022. doi: 10.1038/s41467-021-27801-8.
- [25] S. C. van den Brink, F. Sage, Á. Vértessy, B. Spanjaard, J. Peterson-Maduro, C. S. Baron, C. Robin, and A. van Oudenaarden. Single-cell sequencing reveals dissociation-induced gene expression in tissue subpopulations. *Nature Methods*, 14(10):935–936, 2017. doi: 10.1038/nmeth.4437.

- [26] H. N. Wilkinson and M. J. Hardman. Wound healing: cellular mechanisms and pathological outcomes. *Open Biology*, 10(9):200223, 2020. doi: 10.1098/rsob.200223.
- [27] M. D. Young and S. Behjati. SoupX removes ambient RNA contamination from droplet-based single-cell RNA sequencing data. *GigaScience*, 9(12):giaa151, 2020. doi: 10.1093/gigascience/giaa151.
- [28] K. D. Zimmerman, M. A. Espeland, and C. D. Langefeld. A practical solution to pseudoreplication bias in single-cell studies. *Nature Communications*, 12(1):738, 2021. doi: 10.1038/s41467-021-21038-1.

Supplementary methods and numerical results

S1 Cohort roles and evidence units

Table S1 fixes the role of each accession. Specimen counts, donor counts and cell counts answer different questions. The 54-sample source series is not the 25-sample frozen foot-skin subset, and neither is the 11-patient outcome subset. Cells in an external specimen are not independent replications of that specimen’s clinical label.

Table S1. Cohort roles, analysis units and interpretation limits. GEO accessions identify the public source studies. Counts refer to specimens, mapped patients or donors as explicitly stated, rather than independent cells. DFU denotes diabetic foot ulcer and QC denotes quality control. The original discovery series, selected foot-skin representation and outcome subset have different denominators; external projections test construct associations or feasibility without a qualifying patient-level healing validation.

Cohort	Units	Role and limit
GSE165816	54 original specimens; 25 selected foot specimens	Discovery representation; 14 DFU specimens map to 11 outcome patients. The original series maps to 27 participants.
GSE231643	8 specimens	Title-derived healing labels (5, 3); no authoritative independent patient mapping.
GSE241132	3 donors; 12 timepoint samples	Healthy acute-wound construct audit; no DFU prognosis inference.
GSE223964	8 specimens	Frozen projection and exclusive immune-compartment audit; no qualifying healing outcome.
GSE245703	12 specimens	Frozen projection and immune-compartment audit; no qualifying healing outcome.
GSE268834	8 specimens	Descriptive diabetes-group and immune-context contrasts, not healing validation.
GSE248247	2 specimens; 358 QC cells	Low-yield projection check from 561 raw cells; correlations not estimated.

S2 Operational lineage and mixture rules

Discovery lineage scores average gene-wise standardized $\log(1 + 10^4 x_{ig} / \sum_h x_{ih})$ values within marker sets; the highest lineage score assigns the label. The fibroblast markers are COL1A1, COL1A2, PDGFRA, TWIST2, COL6A3, MMP2 and CCDC80. The circularity check concerns overlap with the tested topic’s specified top-gene list; it does not imply independence from all genes contributing to the neural representation. The loading is a learned function of the full panel.

In external audits, immune assignment is PTPRC detection or detection of at least two markers within a T, NK, myeloid, B or plasma panel. These panels respectively contain CD3D, CD3E, CD2, TRBC1, TRBC2, NKG7, GNLY; LYZ, CD68, CD14, ITGAX, FCER1G, TYROBP, AIF1; MS4A1, CD79A, CD79B, BANK1; and MZB1, JCHAIN, DERL3, TNFRSF17. Fibroblast-like cells require at least two of COL1A1, COL1A2, COL3A1, DCN, LUM, PDGFRA, THY1 and no immune assignment. B/plasma cells are immune-assigned cells detecting at least one B or plasma marker. All remaining cells form the other compartment. Unlike the discovery lineage list, this external structural list includes topic-associated genes; it is a descriptive operational gate and is not evidence of gene-disjoint external validation.

The disjoint immune, fibroblast-like and other compartments sum to one. B/plasma is a subset of immune, so its within-immune share must not be added as a fourth exclusive compartment. The older overlapping marker-detection fractions are not used for the current external associations.

For distributional descriptions, pure high requires the sample’s fibroblast 10th percentile to be at least 0.10; pure low requires its 90th percentile to be at most 0.15; mixed requires a 10th percentile at most 0.05 and a 90th percentile at least 0.40. Remaining samples are labelled other. These operational categories do not prove distinct biological attractors.

S3 Representation integrity is a restricted diagnostic

Let the reference and candidate k -nearest-neighbour graphs have component labels r_i, c_i . If n_{ab} counts cells jointly in reference component a and candidate component b , and $n_a = \sum_b n_{ab}$, the implementation computes reachability preservation exactly as

$$R_{\text{reach}} = \frac{\sum_{a,b} n_{ab}^2 - N}{\sum_a n_a^2 - N}, \quad (\text{S1})$$

Equation (S1) is assigned value one when the denominator is zero. This sparse component-count calculation avoids materializing all pairwise reachability values. It is a connectivity statistic: two connected graphs can give one despite severely distorted local geometry. The separate edge-survival score is reported but is not included in the binary integrity condition.

Mean cell perplexity is $\bar{P} = N^{-1} \sum_i \exp[-\sum_k \theta_{ik} \log \theta_{ik}]$. The recorded integrity condition is $R_{\text{reach}} > 0.70$ and $\bar{P} \leq 0.95K$. Gene-column shuffling raises $\bar{P}$ from 3.68 of 15 to 14.43 of 15 and fails the concentration criterion despite retained connectivity. This rejects the tested near-uniform input; it does not validate every topological feature or every cut at the fixed reachability bound. The K and state-cut grids in the main text do not constitute a sensitivity scan of the 0.70 connectivity bound.

S4 Resampling and future-study assumptions

Every interval names its resampling unit. The semantic bootstrap resamples sample rows. The all-cell outcome bootstrap resamples patients within arms after replicate collapse. AUC intervals resample fixed out-of-fold predictions; they do not refit the encoder or reproduce the entire learning algorithm on each bootstrap draw. Permutation inference reruns outcome-dependent score orientation and representation selection, which is why the appropriate null AUC need not be 0.5. The unadjusted $p = 0.0424$ for the sample-level module is exploratory across several readouts and is not a confirmatory multiplicity-controlled endpoint.

The next-study design calculation uses sample-derived arm spreads. This creates uncertainty not reflected by its Monte Carlo standard error; at estimated power 0.8 and 4,000 simulations, the conditional standard error is about $\sqrt{0.8(0.2)/4000} = 0.0063$. Candidate equivalence designs assume a true difference of zero. Bounded topic scores are approximated by Gaussian locations and variances, and proposed totals near 80% power are approximate search outputs rather than guaranteed minima.

For the exploratory AUC-design simulator, a unit-variance binormal score with class-mean separation d has $\text{AUC} = \Phi(d/\sqrt{2})$, hence $d = \sqrt{2}\Phi^{-1}(\text{AUC})$. Correlation is imposed on within-class noise, not on the full mean-shifted scores. Its standard-error approximation is estimated across simulated AUC differences; it is a model-based planning exercise, not a validated patient-level test. Neither these AUC calculations nor an inpatient design-effect formula licenses treating repeated biopsies as independent prognosis cases.

S5 Numerical results

Tables S2–S9 present the mapped patient scores, discrimination intervals, cross-method associations, conditional design calculations, semantic uncertainty and external immune-context summaries. Their denominators and limits accompany the values so that the three distinct readouts—all-cell mean, fibroblast mean and high-state fraction—remain identifiable. Tables S10 and S11 retain the complete covariate screen and immune/placebo residual summaries. Tables S12, S13, S14 and S15 distinguish the lineage, discrimination, paired-anatomy and equivalence estimands. Tables S18 and S17 retain every tested threshold and topic capacity, including adverse sensitivities. Table S7 records the exploratory AUC-gap simulation. The separate arithmetic conversion from specimens to patients under assumed inpatient correlation is not a validated patient-prognosis design and is not used as a recruitment recommendation.

Table S2. Author-mapped GSE165816 DFU patient scores (7 healed, 4 not healed). Each score is the mean frozen fibroblast-associated topic loading over all retained cells in the listed foot specimens; pooling cells weights repeated specimens by retained cell count. Cells are counts, not independent outcome observations. Patient identifiers are study-derived pseudonyms and G labels identify specimens. The outcome is the source-study category within 12 weeks, not an individual healing day. Scores are on the topic-loading scale from zero to one.

Patient	Outcome	Specimens	Cells	All-cell mean
S01	not healed	G6	2279	0.0500
S02	healed	G23, G4	7495	0.1448
S03	healed	G2	3951	0.0213
S04	healed	G15	2250	0.0749
S05	not healed	G9	4290	0.0364
S06	not healed	G33, G34	7681	0.0131
S07	not healed	G39	2481	0.0305
S08	healed	G7, G8	4192	0.0751
S09	healed	G49	2928	0.0539
S10	healed	G42	1838	0.0201
S11	healed	G45	3374	0.0331

Table S3. Exploratory discrimination across 9 healing and 5 non-healing GSE165816 specimens. AUC is the area under the ROC curve for leave-one-sample-out scores; higher values favour healing. The percentile 95% intervals use 4,000 arm-stratified resamples of fixed out-of-fold scores. One-sided permutation probabilities and null means rerun fold-dependent orientation and fitting. PCA is principal component analysis; NMF is nonnegative matrix factorization. The frozen topic encoder itself used the discovery expression cohort. Repeated patients limit independence, and none of the six paired AUC-difference intervals excludes zero.

Readout	AUC	Bootstrap 95% interval	Permutation p	Null mean
Fibroblast topic	0.822	0.556 to 1.000	0.0604	0.561
Module	0.844	0.578 to 1.000	0.0424	0.497
PCA	0.533	0.156 to 0.889	0.4527	0.480
NMF	0.878	0.655 to 1.000	0.0597	0.553

Table S4. Patient-unit discrimination and anatomical sensitivity within GSE165816. The first four rows are leave-one-patient-out AUCs for 7 healing and 4 non-healing patients, with 4,000 arm-stratified bootstrap draws of fixed scores and one-sided label-permutation probabilities. PCA and NMF denote principal component analysis and nonnegative matrix factorization. The final row is the foot-minus-forearm all-cell fibroblast-associated topic mean difference in 10 paired patients, with a 4,000-draw paired bootstrap interval and a two-sided exhaustive sign-flip probability using add-one correction. Its Benjamini–Hochberg q value across 15 topics is 0.0402. The two row types have different estimands; neither supplies external validation.

Readout	Estimate	Bootstrap 95% interval	Permutation p
Fibroblast topic	0.7143	0.3571 to 1.0000	0.3021
Module	0.7857	0.5000 to 1.0000	0.1662
PCA	0.5000	0.0000 to 1.0000	0.5025
NMF	0.1250	0.0000 to 0.4286	0.9751
foot minus forearm	0.0476	0.0119 to 0.0916	0.0107

Table S5. Spearman correlations between the four representative sample-level prediction scores across the same 14 DFU specimens (9 healing, 5 non-healing). Fibroblast topic is the frozen all-cell loading, Module the specified fibro-inflammatory score, PCA principal component analysis and NMF nonnegative matrix factorization. Each row is one pair of readouts; values are descriptive point estimates without confidence intervals or multiplicity-adjusted tests. Related scores do not establish equivalent discrimination or rank methods.

Readout A	Readout B	Spearman correlation
Fibroblast topic	Module	0.7231
Fibroblast topic	PCA	0.1648
Fibroblast topic	NMF	0.6023
Module	PCA	0.4022
Module	NMF	0.6756
PCA	NMF	0.3200

Table S6. Conditional next-cohort simulations for all-cell fibroblast-associated topic mean loading. Rows assume one independent patient per observation, no missing outcomes, healing allocation near 9/14, and arm SDs 0.04845 and 0.01584 estimated from discovery specimens. Two-sided Welch tests assess candidate mean differences; Welch two-one-sided tests assess equivalence under a true difference of zero, at $\alpha=0.05$. Total and arm columns give the simulated observation counts; power uses 4,000 draws per candidate size. Values near 0.80 have Monte Carlo error and are approximate search outputs. Candidate differences and margins are inputs, not clinical minimum important differences or validated recruitment targets.

Design	Difference / margin	Total	Healing / other	Estimated power
Detect difference	0.02	87	56 / 31	0.802
Detect difference	0.03	41	26 / 15	0.798
Detect difference	0.0435	21	14 / 7	0.822
Detect difference	0.06	12	8 / 4	0.807
Detect difference	0.08	9	6 / 3	0.862
Detect difference	0.1	7	4 / 3	0.786
Equivalence under zero difference	0.02	94	60 / 34	0.805
Equivalence under zero difference	0.03	43	28 / 15	0.803
Equivalence under zero difference	0.05	17	11 / 6	0.807
Equivalence under zero difference	0.075	9	6 / 3	0.829
Equivalence under zero difference	0.1	7	4 / 3	0.856

Table S7. Exploratory binormal planning simulation for paired AUC differences. Two candidate AUCs are centred on 0.80, separated by the stated gap, with within-class score-noise correlation 0.5022 and healing allocation near 9/14. Each searched total uses 4,000 simulation draws and a model-based normal approximation for the paired difference at $\alpha=0.05$. Totals assume one observation per independent patient and no missing outcome. The gaps are hypothetical inputs; the approximate 80% power search does not establish actual clinical performance, calibrated patient-test power or a validated recruitment target.

Candidate AUC gap	Total observations	Healing / non-healing	Estimated power
0.05	721	464 / 257	0.818
0.1	181	116 / 65	0.800
0.15	84	54 / 30	0.816
0.2	52	33 / 19	0.811

Table S8. Association between mean fibroblast-associated topic loading within fibroblasts and the threshold-defined high-state fraction across 25 GSE165816 specimens. The Spearman row gives the observed correlation and percentile 95% limits from 10,000 sample-row bootstrap draws. The leave-one-out row gives the minimum and maximum of 25 recalculations, not a confidence interval; a dash means no single point estimate applies. A separate 10,000-draw permutation analysis gave a two-sided add-one probability of 0.00010 and a central null interval from -0.4025 to 0.4072. Samples may share patients, and the two summaries share cells.

Summary	Estimate	Lower endpoint	Upper endpoint
spearman rho	0.9325	0.7887	0.9828
leave one out rho	—	0.9235	0.9566

Table S9. External sample-level Spearman correlations with the frozen fibroblast-associated topic mean under exclusive lineage gates. Immune fraction divides immune-assigned cells by all retained cells; B/plasma share divides B/plasma cells by immune-assigned cells. Detected genes and unique molecular identifiers (UMIs) are separate per-sample medians over QC cells. All entries are descriptive point estimates without confidence intervals. Specimens are not verified independent patients, and no listed cohort supplies a qualifying patient-level healing outcome.

Cohort	Samples	Immune fraction	B/plasma share	Detected genes	UMIs
GSE223964	8	-0.024	-0.071	-0.429	-0.476
GSE245703	12	0.846	0.077	-0.427	-0.308
GSE268834	8	0.714	-0.500	0.238	0.262

Table S10. Complete exploratory covariate screen in 25 discovery specimens. Each row correlates the fibroblast high-state fraction, using the historical rule loading at least 0.184, with one of 31 covariates. Lineage fractions use all retained cells; technical summaries and topic means use fibroblasts. Values are Spearman rho, two-sided asymptotic p and Benjamini–Hochberg q across all 31 tests, ordered by absolute rho; 14 rows have q below 0.05. UMI means unique molecular identifier. The G label is an identifier, not an established batch or time variable. Repeated patients, shared cells and compositional constraints limit independence; these marginal associations are not causal or healing-validation tests.

Covariate	Specimens	Spearman rho	Raw p	BH q
B/plasma fraction	25	0.8279	3.26e-07	1.01e-05
Fibroblast CFD/APOD mean	25	-0.7664	7.95e-06	0.000123
Fibroblast PTPRC-positive fraction	25	0.7267	3.88e-05	0.000401
Fibroblast median UMI count	25	0.6632	0.000302	0.00234
Fibroblast IGFBP7/TAGLN mean	25	0.6514	0.000421	0.00261
Sweat-gland fraction	25	-0.5832	0.00222	0.0114
Fibroblast FOS/IGFBP7 mean	25	-0.5687	0.00302	0.0134
Fibroblast DCN/CFD mean	25	-0.5592	0.00366	0.0142
Fibroblast IGFBP7/IGFBP5 mean	25	0.5316	0.00624	0.0215
Fibroblast COL1A1/COL1A2 mean	25	0.5238	0.00721	0.0223
Keratinocyte fraction	25	-0.4836	0.0143	0.0404
Fibroblast HLA-DRA/CD74 mean	25	0.4726	0.0171	0.0441
T-cell fraction	25	0.4647	0.0193	0.0459
Fibroblast TAGLN/ACTA2 mean	25	0.4552	0.0222	0.0492
Melanocyte fraction	25	-0.4442	0.0261	0.054
Fibroblast JUNB/DNAJB1 mean	25	0.4182	0.0375	0.0726
Mast-cell fraction	25	-0.3977	0.0489	0.0893
Myeloid fraction	25	0.3891	0.0546	0.094
Numeric G specimen label	25	-0.3607	0.0765	0.125
Endothelial fraction	25	-0.2347	0.259	0.401
Total retained cells	25	-0.2276	0.274	0.404
Fibroblast DNAJB1/FOS mean	25	0.1977	0.344	0.484
Fibroblast KRT14/KRT10 mean	25	0.1457	0.487	0.63
Fibroblast dissociation score	25	0.1449	0.489	0.63
Lymphatic endothelial fraction	25	-0.1331	0.526	0.63
Fibroblast fraction	25	0.1292	0.538	0.63
Fibroblast FOS/JUNB mean	25	-0.1260	0.548	0.63
Retained fibroblasts	25	0.0851	0.686	0.759
Pericyte/smooth-muscle fraction	25	-0.0473	0.823	0.879
Fibroblast IGFBP7/SPARCL1 mean	25	-0.0354	0.866	0.895
Fibroblast DCD/SAT1 mean	25	-0.0236	0.911	0.911

Table S11. Immune-transcript exclusion and residual-mixture controls. BIC denotes Bayesian information criterion, fitted to sample high-state fractions or their regression residuals; the difference is BIC for one Gaussian minus BIC for two. Positive differences favour two components. The first two rows share 24 matched specimens and have fraction correlation 0.9887. Residual regressions use all 25 specimens: immune covariates are B/plasma, T-cell and myeloid fractions plus immune-transcript positivity; placebos are pericyte/smooth-muscle, endothelial, melanocyte and sweat-gland fractions. R squared is variance explained by each regression and is inapplicable to unadjusted rows (dash). Both residual differences are negative, so the analysis does not establish immune-specific separation.

Summary	n	R squared	BIC one	BIC two	Difference
Before immune-transcript exclusion	24	—	12.417	-28.569	40.986
After immune-transcript exclusion	24	—	5.342	-51.341	56.683
Immune-covariate residuals	25	0.4953	-3.900	0.031	-3.931
Placebo-covariate residuals	25	0.3549	2.235	10.433	-8.198

Table S12. Historical sample-level lineage and loading contrasts. Ratios divide the healing-arm mean by the non-healing-arm mean; intervals are percentile 95% limits from 20,000 within-arm sample bootstraps, and p values are two-sided Mann–Whitney tests. GSE165816 uses 9 versus 5 DFU specimens; GSE231643 uses 5 versus 3 title-labelled specimens. The discovery ratios come from the saved lineage-projection analysis and are distinct from the later deterministic all-cell patient contrast. Fibroblast fraction uses all cells as denominator; fibroblast mean uses only assigned fibroblasts. These are separate estimands. Repeated discovery patients and unavailable external patient mappings prevent independent outcome validation.

Cohort	Readout	Mean ratio	Bootstrap 95% interval	Raw p
GSE165816	Fibroblast fraction	1.425	0.873 to 2.298	0.438
GSE165816	All-cell fibroblast-associated topic mean	2.724	1.369 to 5.732	0.0599
GSE165816	Fibroblast fibroblast-associated topic mean	2.889	1.059 to 13.966	0.042
GSE231643	Fibroblast fraction	2.051	1.128 to 4.184	0.0714
GSE231643	All-cell fibroblast-associated topic mean	2.564	1.310 to 8.079	0.143
GSE231643	Fibroblast fibroblast-associated topic mean	1.660	1.007 to 6.102	0.143

Table S13. All six paired AUC differences at each analysis unit. Each contrast subtracts the second leave-one-out readout from the first. Percentile 95% intervals use 4,000 common-index, arm-stratified bootstrap draws of fixed out-of-fold scores: 9 versus 5 specimens or 7 versus 4 mapped patients. PCA and NMF denote principal component analysis and nonnegative matrix factorization. No sample-unit interval excludes zero; patient-unit Fibroblast topic minus NMF and Module minus NMF intervals do. The intervals are exploratory, are not adjusted across pairs, and do not refit the entire representation in each bootstrap draw.

Unit	AUC contrast	Difference	Bootstrap 95% interval
14 specimens	Fibroblast topic minus Module	-0.0222	-0.2889 to 0.2222
14 specimens	Fibroblast topic minus PCA	0.2889	-0.0889 to 0.6667
14 specimens	Fibroblast topic minus NMF	-0.0556	-0.3556 to 0.2222
14 specimens	Module minus PCA	0.3111	0.0000 to 0.6667
14 specimens	Module minus NMF	-0.0333	-0.2889 to 0.2000
14 specimens	PCA minus NMF	-0.3444	-0.7111 to 0.0000
11 patients	Fibroblast topic minus Module	-0.0714	-0.4286 to 0.2857
11 patients	Fibroblast topic minus PCA	0.2143	-0.4643 to 0.9286
11 patients	Fibroblast topic minus NMF	0.5893	0.1429 to 0.9643
11 patients	Module minus PCA	0.2857	-0.1786 to 0.7857
11 patients	Module minus NMF	0.6607	0.2857 to 0.9652
11 patients	PCA minus NMF	0.3750	0.0000 to 0.7857

Table S14. Complete paired anatomical control across 15 topics in 10 author-mapped patients. Differences are mean all-cell loadings in foot minus forearm, after cell-count-weighted collapse of repeated foot specimens. Intervals are percentile 95% limits from 4,000 paired-patient bootstrap draws. Two-sided p values enumerate all 1,024 sign allocations and retain add-one correction; q values use Benjamini–Hochberg across 15 topics. The TIMP1/COL3A1, COL1A1/COL1A2, CFD/APOD and IGFBP7/TAGLN loadings have q below 0.05. Labels list the two leading decoder genes; they do not assert cell-type identity. Bootstrap intervals and permutation tests need not be dual. This is an anatomical sensitivity control within the discovery source, not a healing endpoint.

Leading genes	Foot minus forearm	Bootstrap 95% interval	Sign-flip p	BH q
TIMP1/COL3A1	0.0476	0.0119 to 0.0916	0.0107	0.0402
JUNB/DNAJB1	0.0103	-0.0616 to 0.0793	0.799	0.956
COL1A1/COL1A2	0.0914	0.0535 to 0.1376	0.00293	0.022
IGFBP7/IGFBP5	0.0188	0.0041 to 0.0361	0.0634	0.119
TAGLN/ACTA2	0.0016	-0.0200 to 0.0241	0.893	0.956
KRT14/KRT10	-0.2491	-0.4274 to -0.0682	0.0361	0.0774
CFD/APOD	-0.0362	-0.0567 to -0.0169	0.0107	0.0402
DCN/CFD	-0.0159	-0.0383 to 0.0042	0.208	0.312
DCD/SAT1	-0.0001	-0.0083 to 0.0098	0.975	0.975
IGFBP7/SPARCL1	0.0398	-0.0015 to 0.0897	0.134	0.223
FOS/IGFBP7	-0.0038	-0.0087 to 0.0025	0.245	0.334
HLA-DRA/CD74	0.0111	-0.0591 to 0.0984	0.854	0.956
IGFBP7/TAGLN	0.0715	0.0436 to 0.1039	0.00293	0.022
FOS/JUNB	0.0068	0.0025 to 0.0118	0.0185	0.0556
DNAJB1/FOS	0.0061	0.0020 to 0.0108	0.0224	0.0561

Table S15. Complete Welch two-one-sided-test sensitivity grids for the all-cell mean fibroblast-associated topic contrast. The symmetric margin is on the loading scale, and both lower and upper one-sided p values must be below 0.05 for the final column to be Yes. Sample rows use 9 versus 5 specimens; patient rows use 7 versus 4 mapped patients with cell-count-weighted repeat collapse. Margins were examined retrospectively and have no externally justified clinical meaning. A Yes entry therefore records a mathematical sensitivity result, not demonstrated clinical equivalence.

Unit	Margin	Lower-test p	Upper-test p	Both below 0.05
14 specimens	0.05	0.000143	0.359	No
14 specimens	0.10	3.54e-06	0.00438	Yes
14 specimens	0.20	2.07e-08	1.57e-06	Yes
14 specimens	0.30	6.17e-10	1.22e-08	Yes
14 specimens	0.50	5.18e-12	3.21e-11	Yes
11 patients	0.05	0.00131	0.13	No
11 patients	0.10	5.19e-05	0.00206	Yes
11 patients	0.20	7.01e-07	6e-06	Yes
11 patients	0.30	4.05e-08	1.77e-07	Yes
11 patients	0.50	9.05e-10	2.22e-09	Yes

Table S16. Patient-level robustness in the discovery cohort. The first row correlates fibroblast mean loading with the high-state fraction after weighting repeated specimen summaries by fibroblast count within 20 patients (25 specimens). Its percentile 95% interval uses 10,000 patient resamples, seed 217. The second row uses the previously defined all-cell healing contrast in 7 healing and 4 non-healing patients; its original interval remains in the outcome table. Omission ranges give the minimum and maximum after deleting one patient at a time (20 or 11 omissions), not confidence limits. Neither analysis refits the expression model, changes the state cut, adds patients, or provides external validation.

Quantity	Unit	Estimate	Patient bootstrap 95%	Patient-omission range
Mean versus high-state fraction	20 patients	0.9372	0.7760 to 0.9903	0.9266 to 0.9643
Healing minus non-healing	11 patients	0.0280	Not recalculated	0.0139 to 0.0347

Table S17. Topic-capacity sensitivity after tracking the TIMP1-carrying topic at each K. Gene-based descriptions identify the tracked loading; each fit supplies its own two-Gaussian midpoint cut. Semantic rho uses 25 specimens; the unadjusted healing-arm p value uses 9 versus 5 high-state-fraction summaries. The final column asks whether the maximum fraction occurs in healthy skin. At K=10 the tracked topic is a different IGFBP-rich programme and the counterexample fails. Topics across capacities are not assumed to have identical biological meaning.

K	Programme description	Derived cut	Semantic rho	Healing p	Healthy maximum
10	IGFBP-rich	0.1454	0.9469	0.147	No
15	CHI3L1/COL3A1-rich	0.1691	0.9316	0.109	Yes
20	CHI3L1/COL3A1-rich	0.1695	0.9236	0.109	Yes

Table S18. All 18 state-cut sensitivity evaluations for the frozen K=15 topic. Spearman rho correlates fibroblast mean loading with the threshold-defined high-state fraction across 25 specimens. The ratio compares mean high-state fractions in 9 healing versus 5 non-healing specimens; p is the unadjusted two-sided Mann–Whitney probability. The last column gives the group containing the largest fraction. The grid includes the fitted cut rounded to 0.1843. Four cuts have p below 0.05; these are dependent sensitivity analyses, not 18 independent confirmations.

State cut	Semantic rho	Healing / non-healing	Raw p	Largest fraction
0.0800	0.9326	2.966	0.042	Healthy
0.1000	0.9414	2.959	0.042	Healthy
0.1200	0.9091	3.011	0.042	Healthy
0.1400	0.9301	2.992	0.042	Healthy
0.1600	0.9294	2.981	0.0599	Healthy
0.1800	0.9349	3.001	0.0599	Healthy
0.1843	0.9325	2.975	0.0599	Healthy
0.2000	0.9349	2.987	0.0599	Healthy
0.2200	0.9341	3.124	0.0599	Healthy
0.2400	0.9228	3.352	0.0617	Healthy
0.2600	0.9228	3.424	0.0617	Healthy
0.2800	0.9236	3.465	0.0617	Healthy
0.3000	0.9236	3.499	0.0617	Healthy
0.3200	0.9479	3.741	0.0617	Healthy
0.3400	0.9463	3.699	0.0617	Healthy
0.3600	0.9401	4.040	0.0608	Healthy
0.3800	0.9182	3.945	0.0608	Healthy
0.4000	0.9182	3.889	0.0608	Healthy

Table S19. Exact mean decomposition on the saved 19,410 discovery fibroblasts in 25 specimens. The cut is 0.184295848; empirical bin means are distinct from fitted Gaussian-component means. The constant-state approximation uses the cell-pooled low/high means, and its residual is the observed specimen mean minus that approximation. RMS denotes root-mean-square residual with equal specimen weights. Ranges exclude undefined means for empty bins; zero-weight terms contribute zero to the identity. Both exact identities have maximum numerical discrepancy below 10^{-15} . Values are on the zero-to-one loading scale, apart from the count row. They describe the frozen sample, without confidence intervals, a test of invariant state intensity or causal variance attribution.

Quantity	Value or range	Denominator
Cell-pooled low-bin mean	0.02380	All 19,410 fibroblasts
Cell-pooled high-bin mean	0.51301	All 19,410 fibroblasts
Specimen low-bin means	0.01204 to 0.10213	25 nonempty low bins
Specimen high-bin means	0.22645 to 0.57777	16 nonempty high bins
Specimens with an empty bin	9	9 of 25; high bin empty
Within-state residual range	-0.02434 to 0.04377	All 25 specimens
Within-state residual RMS	0.01575	Equal specimen weights

Table S20. Independent arithmetic check of the all-cell mean-loading contrast in 11 patients (7 healing, 4 non-healing). Exhaustive allocation includes the observed labels and numerical ties: 89/330 allocations are at least as extreme in absolute mean difference, whereas the retained add-one convention gives 90/331. These probabilities assume observational label exchangeability and do not identify a causal effect. Welch quantities use sample variances within the two outcome groups. Standard error, interval and margin are on the loading scale; degrees of freedom and probabilities are dimensionless. A margin strictly greater than the infimum would be needed to reject both one-sided tests at 0.05. The data-derived boundary is not a clinical equivalence margin.

Quantity	Value
All label allocations	330
At least as extreme as observed	89
Exhaustive two-sided probability	0.269697
Recorded add-one probability	0.271903
Welch standard error	0.018226
Welch degrees of freedom	8.098869
Two-sided 90% Welch interval	-0.005864 to 0.061814
Equivalence-margin infimum	0.061814